



TÉCNICO
LISBOA



INTRODUCTION TO NEURAL NETWORKS

Giles Strong

6^as Jornadas de Engenharia Física, IST - 04/03/20

giles.strong@outlook.com

twitter.com/Giles_C_Strong

Amva4newphysics.wordpress.com


github.com/GilesStrong



MOTIVATION FOR NEURAL NETWORKS

FUNCTIONAL APPROXIMATION

- Have data x in some observation space and want to map data to some target space $y: y = f(x)$, e.g.
 - pixel values \rightarrow cat or dog (image classification)
 - particle momenta \rightarrow signal or background (particle collision classification)
 - environment space \rightarrow reward (reinforcement learning)
 - text \rightarrow sentiment (sentiment classification)



$$y = f(x)$$

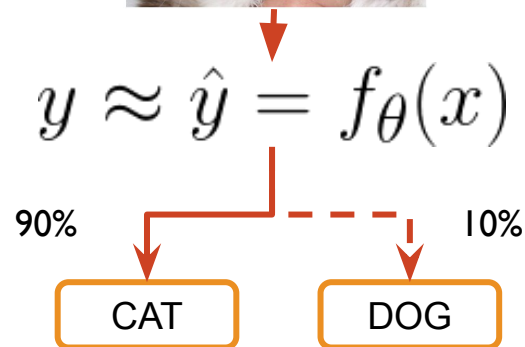

CAT

DOG

FUNCTIONAL APPROXIMATION

- If domain theory doesn't offer a (computationally cheap) analytical map, can instead approximate it with a sufficiently flexible parameterisation:

- $y \approx \hat{y} = f_{\theta}(x)$




LOSS FUNCTIONS

- Aim is to adjust parameters θ in order to make approximated y match true y as closely as possible:
$$\theta = \arg \min_{\theta} \mathcal{L}(y, \hat{y}_{\theta})$$
- Where L is a function which quantifies the error (difference between prediction and target): the *loss function*
- Exact form of loss function varies from problem to problem
- However, making f_{θ} sufficiently flexible requires many continuously valued parameters: finding the optimal parameters by brute force is intractable


PARAMETER OPTIMISATION

- A neural network is special form of f_{θ} in which the prediction is continuously differentiable with respect to the parameters
 - [Interactive demo](#)



PARAMETER OPTIMISATION


- By also using a continuously differentiable loss function, the exact effect of each parameter on the performance of the network can be analytically computed:
 - $\nabla_{\theta_k} \mathcal{L}$
- This is done via *backpropagation*
- See e.g. [Stanford CS231n](#) for an in-depth explanation



PARAMETER OPTIMISATION

- The optimal parameters can now be converged to via iterative updates in the direction of steepest slope of the loss function:

$$\theta_{k+1} = \theta_k - \gamma \nabla_{\theta_k} \mathcal{L}$$



PARAMETRIC FORM

- Parameters (mostly) take the form of multiplicative matrices (weights) and additive offsets (biases) and are said to be arranged in *layers*, each of which apply a linear transformation to the previous layer:
 - $w \cdot x + b$
- By interspersing these linear transformations with non-linear *activation functions* more complex transformations can be achieved:
 - $\mathcal{A}(w \cdot x + b)$
- Stacking linear layers and activation functions produces e.g. a 2-layer network:
 - $\mathcal{A}(w_1 \cdot \mathcal{A}(w_0 \cdot x + b_0) + b_1)$

PARAMETRIC FORM

- Explicitly, the linear transformation is:

$$\begin{pmatrix} w_{0,0} & w_{0,1} & \dots & w_{0,N} \\ w_{1,0} & & & \\ & & & \\ \dots & & & \\ w_{M,0} & & & w_{M,N} \end{pmatrix} \times \begin{bmatrix} x_0 \\ x_1 \\ \dots \\ x_N \end{bmatrix} + \begin{bmatrix} b_0 \\ b_1 \\ \dots \\ b_M \end{bmatrix} = \begin{bmatrix} h_0 \\ h_1 \\ \dots \\ h \end{bmatrix}$$


- And the activation is then applied to each element of the resulting vector via broadcasting:
 - $\mathcal{A}(\bar{h})$



COMPONENTS OF NEURAL NETWORKS


ACTIVATION FUNCTIONS

- Activation functions allow the network to perform non-linear transformations
 - Can learn to model complex functions more easily
- Any function can be used, provided it is continuously differentiable
 - Need to propagate gradients
- Modern standard is ReLU
 - But new ones are being developed
 - E.g. SELU and Swish



ACTIVATION FUNCTIONS


- The final activation function is the *output activation function*
 - It determines the range of outputs of the network
- The form used is therefore determined by the problem being solved, e.g.:
 - Binary Classification: Sigmoid
 - Range $[0, 1]$
 - Multi-class classification: Softmax
 - Range $[0, 1]$ and output sum = 1
 - Multi-label classification: Sigmoid
 - Regression: linear
 - Range $[-\infty, \infty]$




Sigmoid $\sigma(x) = \frac{1}{1 + e^{-x}}$

PARAMETER INITIALISATION

- The parameters must begin from some random starting values
- If these starting values are of the wrong scale, the network can be difficult to train
 - Interactive demo: <https://www.deeplearning.ai/ai-notes/initialization/>
- Several *initialisation schemes* exist, depending on the activation function used:
 - Sigmoid, linear: [Glorot/Xavier](#)
 - ReLU: [Kaiming/He](#)
 - [LSUV](#) initialisation runs a test loop using training data to rescale starting weights appropriately




Activations for
Glorot init




Activations for
Uniform init

PREPROCESSING

- These initialisation schemes generally aim to output a unit-Gaussian distribution when the input is unit-Gaussian
- This means that for optimal training, the input data should have a mean of zero and a standard deviation of one
- Transforming the input data is called *preprocessing*
- In this case we should subtract the mean of each feature for the training data and divide by their standard deviations



Original feature



Preprocessed feature

LOSS FUNCTIONS

- The loss function defines the problem to be solved by quantifying what is a 'good' and 'bad' prediction
- Several standard ones exist, e.g.:
 - Binary Classification: binary cross-entropy
 - Multi-class classification: categorical cross-entropy
 - Multi-label classification: binary cross-entropy
 - Regression: Mean squared-error or mean absolute-error
- Loss functions can also be *weighted* to account for data-points differently
 - E.g. to encourage unbiased responses


Average over data points

Difference between prediction and truth

$$MSE = \frac{1}{N} \sum_{n=1}^N (\hat{y}_n - y_n)^2$$

OPTIMISERS

- The default optimiser updates in the direction of steepest descent of the loss function by a fixed amount
- But potentially $O(100,000)$ parameters = $O(100,000)$ dimensional space
- Very likely to end up in “narrow valleys”, where steepest descent results in oscillation between valley walls




How GD moves

Ideal moves

OPTIMISERS

- Various improvements to *vanilla* gradient descent:
 - Accumulate *momentum*
 - Cancels out in wall-to-wall direction
 - Accumulates in valley floor direction
 - Scale learning rate (step-size) per parameter based on past gradients
 - Smaller steps on steeper slopes
 - Larger steps on shallower slopes
 - [Interactive demo](#)
 - [Overview](#)
- [Adam](#) is generally a good starting point





TUTORIAL







PYTORCH TUTORIAL

- https://github.com/GilesStrong/PyTorch_Tutorial
- Jupyter notebooks runnable on Google Colab

Click the badges to
load in Colab

Overview

This tutorial is designed to present neural networks from a practical perspective. Several notebooks are found in the `notebooks` folder.

1.  [Open in Colab](#) SGD_from_scratch uses a single neuron implementing backpropagation and weight updates manually.
2.  [Open in Colab](#) Basic_Classification introduces PyTorch's `nn.Module`.
3.  [Open in Colab](#) Basic_Regression again uses PyTorch's `nn.Module`.
4.  [Open in Colab](#) 3_Classification_Application is an example of a simple classification application.
5.  [Open in Colab](#) 4_Regression_Application_Exercise is a regression application.
6.  [Open in Colab](#) 5_Regression_Application_Completed is a more advanced regression application, including a few tricks for training a regressor.

GOOGLE COLAB RUNTIME

0_SGD_from_scratch.ipynb

File Edit View Insert Runtime Tools Help Last saved at 10:44 AM

+ Code + Text

Basic NN exam

In this notebook we'll ex

distribution) and classifi

Neural networks are ma

Runtime menu options:

- Run all %/Ctrl+F9
- Run before %/Ctrl+F8
- Run the focused cell %/Ctrl+Enter
- Run selection %/Ctrl+Shift+Enter
- Run after %/Ctrl+F10
- Interrupt execution %/Ctrl+M
- Restart runtime... %/Ctrl+M
- Restart and run all...
- Factory reset runtime
- Change runtime type**
- Manage sessions
- View runtime logs

Notebook settings

Runtime type
Python 3

Hardware accelerator
None

☐ Omit code cell output when saving this notebook

CANCEL SAVE

Notebook settings

Runtime type
Python 3

Hardware accelerator
GPU

☐ Omit code cell output when saving this notebook

CANCEL SAVE

WORKING IN COLAB

- Colab is similar to Jupyter notebooks
- Run cells by clicking in them and pressing:
 - Control + Enter to run the cell and keep it selected
 - Shift + Enter to run the cell and select the next
- “Runtime” → “Restart Runtime” used to restart and clear memory
- “Runtime” → “Manage Sessions” used to shut down old sessions



FURTHER LEARNING RESOURCES

LIBRARIES

- Most ML development done in Python 3
- Two main libraries: [PyTorch](#) & [TensorFlow](#)
- Both relatively low-level = need good understanding of NNs to use directly; but wrapper libraries exist to provide high-level APIs, e.g.
 - [Keras](#) - no longer developed standalone, but now included in TensorFlow 2.x
 - [Fast.AI](#) - PyTorch wrapper with best practices for image, text, & tabular data but doesn't support weighted data
 - [LUMIN](#) - My own library (in beta) - PyTorch wrapper with best practices for weighted tabular data, plus utilities for HEP, statistics, and interpretation

THEORY & PRACTICE: COURSES

- Fast.AI - free, practical courses; videos + library; top-down experiment first, theory later teaching style:
 - [Machine learning](#) - Fundamentals for data science + Python programming
 - [Deep learning I](#) - Best practices for image, text, & tabular data
 - [Deep learning II](#) - Building DNNs from scratch
- [Stanford course](#) - YouTube lecture series on theory of NNs
- [Yandex MLHEP course](#) - annual week-long intensive introduction to ML for HEP

THEORY & PRACTICE: EXPERIENCE

- Kaggle - data science challenge platform; wide range of challenges, get to see how others approach problems
- Paper reimplementation - helps get more familiar with library, and comfortable changing parts of it, e.g. SELU activation, categorical embedding, learning-rate annealing, and weight averaging

INTERESTING & USEFUL PAPERS

- A disciplined approach to neural network hyper-parameters: Part I -- learning rate, batch size, momentum, and weight decay - [Smith 2018](#)
- SGDR: Stochastic Gradient Descent with Warm Restarts - [Loshchilov & Hutter, 2016](#)
- Entity Embeddings of Categorical Variables - [Guo & Berkhahn, 2016](#)
- Regularization for Deep Learning: A Taxonomy - [Kukačka, Golkov, & Cremers, 2017](#)